

Lecture 3: Data Manipulation with dplyr

Harris Coding Camp – Accelerated Track

Summer 2022

dplyr

Data manipulation with dplyr

The dplyr library provides a toolkit for data manipulation.

We will cover:

- ▶ `select()` to pick columns
- ▶ `filter()` to get rows that meet a criteria
- ▶ `arrange()` to order the data
- ▶ `mutate()` to create new columns
- ▶ `summarize()` to summarize data

tidyverse origins: dplyr

```
library(tidyverse)
## -- Attaching packages -----
## v ggplot2 3.3.0   v purrr   0.3.4
## v tibble  2.1.3   v dplyr   0.8.5
## v tidyr   1.0.2   v stringr 1.4.0
## v readr   1.3.1   v forcats 0.5.0
## -- Conflicts -----
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()    masks stats::lag()
```

dplyr concepts

For each function:

- ▶ data in the first position
- ▶ ... in the second position (i.e. allows for many changes at once)

As I show you examples, I'll work with variations of `txhousing`, `txhousing_short`, `txhousing_narrow` and `txhousing_example`

selecting columns with select()

select()

storms

storm	wind	pressure	date
Alberto	110	1007	2000-08-12
Alex	45	1009	1998-07-30
Allison	65	1005	1995-06-04
Ana	40	1013	1997-07-01
Arlene	50	1010	1999-06-13
Arthur	45	1010	1996-06-21



storm	pressure
Alberto	1007
Alex	1009
Allison	1005
Ana	1013
Arlene	1010
Arthur	1010

selecting columns with select()

Use case: You want to present a subset of your **columns**

```
select(txhousing_short, city, date, sales, listings)
```

```
## # A tibble: 6 x 4
##   city      date sales listings
##   <chr>   <dbl> <dbl>     <dbl>
## 1 Abilene 2000      72      701
## 2 Abilene 2000.     98      746
## 3 Abilene 2000.    130      784
## 4 Abilene 2000.     98      785
## 5 Abilene 2000.    141      794
## 6 Abilene 2000.    156      780
```

think: `select()` extends `[, col_expressions]`

`select` provides a suite of short cuts for selecting columns.

- ▶ Notice that `select` can operate with column names while `[` requires Characters.

```
identical(  
  select(txhousing, city, date, sales, listings),  
  txhousing[c("city", "date", "sales", "listings")],  
)
```

```
## [1] TRUE
```


selecting columns with `select()`

As with `[`, `-` says to exclude the columns listed in the vector.

```
select(txhousing_short, -c(city, date, sales, listings))
```

selecting columns with select()

Use case: I want a bunch of columns with similar names.

- ▶ There are several useful functions, see `?tidyselect::select_helpers`.
- ▶ (For more information see r4ds chapter 5.4)

```
# BaseR requires more coding knowledge  
# txhousing[,grep("^city", names(txhousing))]
```

```
select(txhousing_short, starts_with("city"))
```

```
## # A tibble: 6 x 1
```

```
##   city
```

```
##   <chr>
```

```
## 1 Abilene
```

```
## 2 Abilene
```

```
## 3 Abilene
```

```
## 4 Abilene
```

```
## 5 Abilene
```

selecting columns with select(), helpers

Use case: You want to reorder your columns

- ▶ Notice we used a “select_helpers” function everything().

```
select(txhousing_short,  
       date, everything())
```

```
## # A tibble: 6 x 9
```

```
##   date city      year month sales  volume median listing  
##   <dbl> <chr>   <int> <int> <dbl>    <dbl> <dbl>    <dbl>  
## 1 2000  Abilene  2000     1     72 5380000  71400      7  
## 2 2000. Abilene  2000     2     98 6505000  58700      7  
## 3 2000. Abilene  2000     3    130 9285000  58100      7  
## 4 2000. Abilene  2000     4     98 9730000  68600      7  
## 5 2000. Abilene  2000     5    141 10590000 67300      7  
## 6 2000. Abilene  2000     6    156 13910000 66900      7
```

choose rows that match a condition with `filter()`

`filter()`

storms

storm	wind	pressure	date
Alberto	110	1007	2000-08-12
Alex	45	1009	1998-07-30
Allison	65	1005	1995-06-04
Ana	40	1013	1997-07-01
Arlene	50	1010	1999-06-13
Arthur	45	1010	1996-06-21



storm	wind	pressure	date
Alberto	110	1007	2000-08-12
Ana	40	1013	1997-07-01

choose rows that match a condition with filter()

Get all the data from 2013

```
filter(txhousing, year == 2013)
```

```
## # A tibble: 552 x 9
```

```
##   city      year month sales    volume median listings in
```

```
##   <chr>    <int> <int> <dbl>    <dbl>    <dbl>    <dbl>
```

```
## 1 Abilene  2013     1   114 15794494 125300     966
```

```
## 2 Abilene  2013     2   140 16552641  94400     943
```

```
## 3 Abilene  2013     3   164 19609711 102500     958
```

```
## 4 Abilene  2013     4   213 27261796 113700     948
```

```
## 5 Abilene  2013     5   225 31901380 130000     923
```

```
## 6 Abilene  2013     6   209 29454125 127300     960
```

```
## 7 Abilene  2013     7   218 32547446 140000     969
```

```
## 8 Abilene  2013     8   236 30777727 120000     976
```

```
## 9 Abilene  2013     9   195 26237106 127500     985
```

```
## 10 Abilene 2013    10   167 21781187 119000     993
```

```
## # ... with 542 more rows
```

think: `filter()` extends `[row_expression,]`

```
identical(  
  filter(txhousing, year == 2013),  
  txhousing[txhousing$year == 2013, ]  
)
```

```
## [1] TRUE
```

Notice that `filter` can operate with column names while `[` requires that you use a vector.

filter() drops comparisons that result in NA

Compare:

```
df <- tibble(x = c(1, 2, NA))  
filter(df, x > 1)
```

```
## # A tibble: 1 x 1  
##       x  
##   <dbl>  
## 1     2
```

```
df[df$x > 1, ]
```

```
## # A tibble: 2 x 1  
##       x  
##   <dbl>  
## 1     2  
## 2    NA
```

Recall: Relational operators return TRUE or FALSE

Before moving forward with `filter()`, let's review relational operators and logical operators

Operator	Name
<code><</code>	less than
<code>></code>	greater than
<code><=</code>	less than or equal to
<code>>=</code>	greater than or equal to
<code>==</code>	equal to
<code>!=</code>	not equal to
<code>%in%</code>	matches something in

Recall: Relational operators in practice

```
4 < 4
```

```
## [1] FALSE
```

```
4 >= 4
```

```
## [1] TRUE
```

```
4 == 4
```

```
## [1] TRUE
```

```
4 != 4
```

```
## [1] FALSE
```

```
4 %in% c(1, 2, 3)
```

```
## [1] FALSE
```

%in% operator

Tests if left-hand side is in right-hand side.

```
# 1 %in% c(1, 2, 3, 4)  TRUE
# 5 %in% c(1, 2, 3, 4)  FALSE
c(1, 5) %in% c(1, 2, 3, 4)
```

```
## [1]  TRUE FALSE
```

It operates *element-by-element* on the left-hand side!

Review: logical operators combine TRUEs and FALSEs logically

Operator	Name
!	not
&	and
	or

```
# not true  
! TRUE
```

```
## [1] FALSE
```

```
# are both x & y TRUE?  
TRUE & FALSE
```

```
## [1] FALSE
```

```
# is either x | y TRUE?  
TRUE | FALSE
```

```
## [1] TRUE
```

Review: What do the following return?

```
! (4 > 3)  
(5 > 1) & (5 > 2)  
(4 > 10) | (20 > 3)
```

Review: What do the following return?

Logical operators team up with relational operators.

- ▶ First, evaluate the relational operator
- ▶ Then, carry out the logic.

```
! (4 > 3)           # ! TRUE
(5 > 1) & (5 > 2)    # TRUE & TRUE
(4 > 10) | (20 > 3)  # FALSE | TRUE
```

Recall: What do the following return?

```
! (4 > 3)           # ! TRUE
```

```
## [1] FALSE
```

```
(5 > 1) & (5 > 2)    # TRUE & TRUE
```

```
## [1] TRUE
```

```
(4 > 10) | (20 > 3)  # FALSE | TRUE
```

```
## [1] TRUE
```

choose rows that match a condition with `filter()`

Get all the data from 2013 and beyond for Houston.

- ▶ in `filter()` additional match criteria are treated like and

```
filter(txhousing_narrow,  
       year >= 2013,  
       city == "Houston")
```

```
## # A tibble: 3 x 2  
##   city      year  
##   <chr>   <int>  
## 1 Houston  2013  
## 2 Houston  2014  
## 3 Houston  2015
```

choose rows that match a condition with filter()

To do the same operation with [...

```
identical(  
  filter(txhousing,  
    year >= 2013,  
    city == "Houston"),  
  
txhousing[txhousing$year >= 2013 &  
          txhousing$city == "Houston", ]  
)
```

```
## [1] TRUE
```


Why do we get 0 rows here?

Get all the data from 2013 and beyond for Houston and Austin

```
filter(txhousing_narrow,  
       year >= 2013,  
       city == "Houston",  
       city == "Austin")
```

```
## # A tibble: 0 x 2
```

```
## # ... with 2 variables: city <chr>, year <int>
```

choose rows that match a condition with filter()

There's no rows where city is both Houston AND Austin!

We logically want data from Houston OR Austin

```
filter(txhousing_narrow,  
       year >= 2013,  
       city == "Houston" | city == "Austin")
```

```
## # A tibble: 6 x 2  
##   city      year  
##   <chr>    <int>  
## 1 Austin    2013  
## 2 Austin    2014  
## 3 Austin    2015  
## 4 Houston   2013  
## 5 Houston   2014  
## 6 Houston   2015
```

choose rows that match a condition with filter()

At some point you will make this mistake!

```
filter(txhousing_narrow,  
       year >= 2013,  
       city == "Houston" | "Austin")
```

```
Error in filter(txhousing, year >= 2013, city == "Houston"  
Caused by error in 'city == "Houston" | "Austin"':  
! operations are possible only for numeric, logical or comp
```

choose rows that match a condition with `filter()`

What if we want data from Houston, Austin OR Galveston

```
filter(txhousing,  
       year >= 2013,  
       city == "Houston" | city == "Austin" | city == "Ga
```

There has to be an easier way!

choose rows that match a condition with `filter()`

Use `%in%`!

```
in_three_cities <-  
  filter(txhousing,  
         year >= 2013,  
         city %in% c("Houston", "Dallas", "Austin"))
```

Warning: == with two vectors is a bad idea.

Why does it fail to produce the same result?

```
eq_three_cities <-  
  filter(txhousing,  
         year >= 2013,  
         city == c("Houston", "Dallas", "Austin"))
```

```
## Warning in city == c("Houston", "Dallas", "Austin"): longer  
## a multiple of shorter object length
```

```
identical(in_three_cities, eq_three_cities)
```

```
## [1] FALSE
```

```
nrow(in_three_cities)
```

```
## [1] 93
```

```
nrow(eq_three_cities)
```

Warning: == with two vectors is a bad idea.

Why not? Vector recycling.

```
ex <- tibble(id = 1:4,  
             attribute = c("a", "a", "b", "b"))  
  
filter(ex,  
       attribute == c("a", "c"))
```

```
## # A tibble: 1 x 2  
##       id attribute  
##   <int> <chr>  
## 1     1 a
```

Warning: `==` with two vectors is a bad idea.

Why not? Vector recycling.

```
# a == a  
# a == c  
# b == a  
# b == c
```

```
filter(ex, attribute == c("a", "c"))
```

```
## # A tibble: 1 x 2  
##       id attribute  
##   <int> <chr>  
## 1     1 a
```


Crisis averted

```
# a %in% c(a, c)
# a %in% c(a, c)
# b %in% c(a, c)
# b %in% c(a, c)
```

```
filter(ex, attribute %in% c("a", "c"))
```

```
## # A tibble: 2 x 2
##       id attribute
##   <int> <chr>
## 1     1     a
## 2     2     a
```

Introducing the pipe operator



Interlude: Ceci est une %>%

The pipe %>% operator takes the left-hand side and makes it *input* in the right-hand side.

- ▶ by default, the left-hand side is the *first argument* of the right-hand side function.

```
# a tibble is the first argument
select(txhousing, city, year, sales, volume)

txhousing %>%
  select(city, year, sales, volume)
```

Ceci est une %>%

We can chain together tidyverse functions to avoid making so many intermediate data frames!

- ▶ Coming up with names is hard.
- ▶ Updating an object repeatedly leads to a buggy development process

```
txhousing %>%  
  select(city, year, month, median) %>%  
  filter(year == 2013) %>%  
  head()
```

```
## # A tibble: 6 x 4  
##   city      year month median  
##   <chr>   <int> <int>   <dbl>  
## 1 Abilene  2013     1 125300  
## 2 Abilene  2013     2  94400  
## 3 Abilene  2013     3 102500  
## 4 Abilene  2013     4 113700
```

Ceci est une %>%

Read %>% as “and then.

```
# Take data  
txhousing %>%  
  # And then select city, year, month and median  
  select(city, year, month, median) %>%  
  # And then filter where year is 2013  
  filter(year == 2013) %>%  
  # And then show the head (i.e. first 6 rows)  
  head()
```

sort rows with arrange()

arrange()

storms

storm	wind	pressure	date
Alberto	110	1007	2000-08-12
Alex	45	1009	1998-07-30
Allison	65	1005	1995-06-04
Ana	40	1013	1997-07-01
Arlene	50	1010	1999-06-13
Arthur	45	1010	1996-06-21



storm	wind	pressure	date
Ana	40	1013	1997-07-01
Alex	45	1009	1998-07-30
Arthur	45	1010	1996-06-21
Arlene	50	1010	1999-06-13
Allison	65	1005	1995-06-04
Alberto	110	1007	2000-08-12

sort rows with arrange()

```
identical(  
  # base R  
  txhousing[order(txhousing$year), ],  
  
  # base R  
  arrange(txhousing, year)  
)
```

```
## [1] TRUE
```

sort rows with arrange()

To sort in desc()ending order.

```
identical(  
  # base R  
  txhousing[order(txhousing$year,  
                  decreasing = TRUE), ],  
  
  # base R  
  txhousing %>% arrange(desc(year))  
)
```

```
## [1] TRUE
```


creating columns with mutate()

mutate()

storm	wind	pressure	date
Alberto	110	1007	2000-08-12
Alex	45	1009	1998-07-30
Allison	65	1005	1995-06-04
Ana	40	1013	1997-07-01
Arlene	50	1010	1999-06-13
Arthur	45	1010	1996-06-21



storm	wind	pressure	date	ratio	inverse
Alberto	110	1007	2000-08-12	9.15	0.11
Alex	45	1009	1998-07-30	22.42	0.04
Allison	65	1005	1995-06-04	15.46	0.06
Ana	40	1013	1997-07-01	25.32	0.04
Arlene	50	1010	1999-06-13	20.20	0.05
Arthur	45	1010	1996-06-21	22.44	0.04

creating columns in Base R

We've hinted at how to add columns in base R.

What would you guess based on how we added values to lists or changed values in vectors?

creating columns in Base R

1. Refer to a column with `$`.¹
2. Assign a vector to the name

```
# convert the dollar amount to millions  
txhousing$volume_millions <-  
  round(txhousing$volume / 1000000, 1)
```

¹or `[[`

creating columns in Base R

Since we used an assignment operator. The change is permanent.

```
txhousing %>%  
  select(starts_with("volume")) %>%  
  head()
```

```
## # A tibble: 6 x 2  
##       volume volume_millions  
##       <dbl>         <dbl>  
## 1  5380000          5.4  
## 2  6505000          6.5  
## 3  9285000          9.3  
## 4  9730000          9.7  
## 5 10590000         10.6  
## 6 13910000         13.9
```

All my x's live in Texas

Vectors of size 1 are recycled.

```
txhousing$state <- "Texas"
```

```
head(txhousing[, c("city", "state")])
```

```
## # A tibble: 6 x 2
```

```
##   city      state
```

```
##   <chr>    <chr>
```

```
## 1 Abilene Texas
```

```
## 2 Abilene Texas
```

```
## 3 Abilene Texas
```

```
## 4 Abilene Texas
```

```
## 5 Abilene Texas
```

```
## 6 Abilene Texas
```

creating columns with mutate()

mutate works like other dplyr verbs

```
# the data in the first position
mutate(txhousing,
       # after that ... create columns like so
       volume_millions = volume / 1e6)

# same idea
txhousing$volume_millions <-
  txhousing$volume / 1e6
```

- ▶ As with other dplyr verbs, we can refer to the columns by name

creating columns with mutate()

```
identical(  
  # BAD: extracting the column (not as nice)  
  txhousing_example %>%  
    mutate(volume_millions = txhousing_example$volume / 1e6)  
  
  # GOOD: referring to the column by name!  
  txhousing_example %>%  
    mutate(volume_millions = volume / 1e6)  
)
```

```
## [1] TRUE
```

creating columns with mutate()

```
txhousing_example %>%  
  mutate(mean_price = volume / sales) %>%  
  head()
```

```
## # A tibble: 6 x 6  
##   city      year month sales   volume mean_price  
##   <chr>   <int> <int> <dbl>   <dbl>     <dbl>  
## 1 Abilene  2000     1     72  5380000  74722.  
## 2 Abilene  2000     2     98  6505000  66378.  
## 3 Abilene  2000     3    130  9285000  71423.  
## 4 Abilene  2000     4     98  9730000  99286.  
## 5 Abilene  2000     5    141 10590000  75106.  
## 6 Abilene  2000     6    156 13910000  89167.
```


creating columns with mutate()

When we mutate, you can create new columns.

- ▶ *Right-hand side*: the name of a new column.
- ▶ *Left-hand side*: code that creates a vector

```
txhousing_example %>%  
  mutate(mean_price = volume / sales) %>%  
  head()
```

```
## # A tibble: 6 x 6  
##   city      year month sales  volume mean_price  
##   <chr>   <int> <int> <dbl>    <dbl>      <dbl>  
## 1 Abilene  2000     1     72  5380000    74722.  
## 2 Abilene  2000     2     98  6505000    66378.  
## 3 Abilene  2000     3    130  9285000    71423.  
## 4 Abilene  2000     4     98  9730000    99286.  
## 5 Abilene  2000     5    141 10590000    75106.  
## 6 Abilene  2000     6    156 13910000    89167.
```

creating columns with mutate()

You can create multiple columns at a single time and even use information from a newly created column as input.

```
txhousing_example %>%  
  mutate(mean_price = volume / sales,  
         sqrt_mean_price = sqrt(mean_price)) %>%  
  head()
```

```
## # A tibble: 6 x 7
```

```
##   city      year month sales   volume mean_price sqrt_mean  
##   <chr>   <int> <int> <dbl>    <dbl>      <dbl>  
## 1 Abilene  2000     1     72  5380000    74722.  
## 2 Abilene  2000     2     98  6505000    66378.  
## 3 Abilene  2000     3    130  9285000    71423.  
## 4 Abilene  2000     4     98  9730000    99286.  
## 5 Abilene  2000     5    141 10590000    75106.  
## 6 Abilene  2000     6    156 13910000    89167.
```

creating columns with mutate()

Of course, if you want to change the data, you need to assign the output to a name!

```
texas_housing_again <-  
txhousing_example %>%  
  mutate(mean_price = volume / sales,  
         sqrt_mean_price = sqrt(mean_price))
```

Pop quiz

If you load `tidyverse`, you can access the `midwest` data

What `dplyr` function would you need to ...

- ▶ choose the columns `county`, `state`, `poptotal`, `popdensity`
- ▶ get the counties with population over a million
- ▶ reorder the columns by population total
- ▶ round the `popdensity` to the nearest whole number

You try it

- ▶ `select()` the columns `county`, `state`, `poptotal`, `popdensity`
- ▶ `filter()` the counties with population over a million
- ▶ `arrange()` the columns by population total
- ▶ `mutate()` to round the `popdensity` to the nearest whole number
- ▶ `AND mutate()` to round the population totals to the nearest 1000

if you finish early: Try to write it in base R

```
## # A tibble: 4 x 4
##   county    state poptotal popdensity
##   <chr>    <chr>    <dbl>      <dbl>
## 1 COOK      IL      5105000    88018
## 2 WAYNE     MI      2112000    60334
## 3 CUYAHOGA OH      1412000    54313
## 4 OAKLAND   MI      1084000    19702
```

solution

```
midwest %>%  
  select(county, state, poptotal, popdensity) %>%  
  filter(poptotal > 1e6) %>%  
  arrange(desc(poptotal)) %>%  
  mutate(popdensity = round(popdensity),  
         poptotal = poptotal - poptotal %% 1000  
  # alternatively:  
  #   poptotal = round(poptotal, -3)  
  #   poptotal = poptotal - poptotal %% 1000  
  )
```

solution in base R

```
out <- midwest[midwest$poptotal > 1e6,  
               c("county", "state", "poptotal", "popdensity")]  
  
out$popdensity <- round(out$popdensity)  
out$poptotal <- out$poptotal - out$poptotal %% 1000  
  
out[order(out$poptotal, decreasing = TRUE), ]
```

```
## # A tibble: 4 x 4  
##   county    state poptotal popdensity  
##   <chr>    <chr>    <dbl>      <dbl>  
## 1 COOK      IL      5105000    88018  
## 2 WAYNE     MI      2111000    60334  
## 3 CUYAHOGA OH      1412000    54313  
## 4 OAKLAND   MI      1083000    19702
```

solution in base R

- ▶ You may also see some base R functions that are antecedents to dplyr
- ▶ (comparison of subset and filter on stackoverflow)[<https://stackoverflow.com/questions/39882463/difference-between-subset-and-filter-from-dplyr>]

```
# subset: filters and selects referring to names
```

```
out <- subset(midwest, poptotal > 1e6,  
              c(county, state, poptotal, popdensity))
```

```
# within: mutate but with unfortunate syntax
```

```
out <- within(out,{  
              popdensity <- round(popdensity)  
              poptotal <- poptotal - poptotal %% 1000  
            })
```

```
out[order(out$poptotal, decreasing = TRUE), ]
```


summarize data with summarize()

city	particle size	amount ($\mu\text{g}/\text{m}^3$)
New York	large	23
New York	small	14
London	large	22
London	small	16
Beijing	large	121
Beijing	small	56



median
22.5

summarize data with summarize()

Calculate total volume of sales in Texas from 2014.

```
txhousing %>%  
  filter(year == 2014) %>%  
  summarize(total_volume = sum(volume))
```

```
## # A tibble: 1 x 1  
##   total_volume  
##           <dbl>  
## 1  84760948831
```

summarize data with summarize()

Calculate the mean and median number of sales in Texas's three largest cities.

```
txhousing %>%  
  filter(city %in%  
          c("Houston", "Dallas", "San Antonio")) %>%  
  summarize(median_n_sales = median(sales),  
            mean_n_sales = mean(sales))
```

```
## # A tibble: 1 x 2  
##   median_n_sales mean_n_sales  
##           <dbl>         <dbl>  
## 1           3996         3890.
```

summarize data with summarize()

There are many useful functions that go with summarize. Try ?summarize for more.

```
txhousing %>%  
  summarize(n_obs = n(),  
            n_cities = n_distinct(city))
```

```
## # A tibble: 1 x 2  
##   n_obs n_cities  
##   <int>   <int>  
## 1  8602     46
```

Alert: summarize() without summarizing

Weird behavior:

```
# in older versions of dplyr this gives an error  
# Error: Column `mean_price` must be length 1 (a summary value)  
  
txhousing %>%  
  summarize(mean_price = volume / sales) %>%  
  head()
```

piping dplyr verbs together

dplyr verbs can be piped together in any order you want, although different orders can give you different results, so be careful!

```
txhousing %>%  
  select(city, year, month, sales, volume) %>%  
  mutate(log_mean_price = log(volume / sales)) %>%  
  filter(year == 2013) %>%  
  summarize(log_mean_price_2013 = mean(log_mean_price,  
                                         na.rm = TRUE))
```

Won't give you the same result as

txhousing %>%

select(city, year, month, sales, volume) %>%

mutate(log_mean_price = log(volume / sales)) %>%

summarize(log_mean_price = mean(log_mean_price, na.rm = TRUE)) %>%

filter(year == 2013)

Actually this code will give you an error, try it!

Recap: manipulating data with dplyr

We learned

- ▶ how to employ the 5 dplyr verbs of highest importance including
 - ▶ `select()` to pick columns
 - ▶ `arrange()` to order the data
 - ▶ `mutate()` to create new columns
 - ▶ `filter()` to get rows that meet a criteria
 - ▶ `summarize()` to summarize data
- ▶ how to use relation operators, binary operators for math and logical operators in dplyr contexts

Next steps:

Lab:

- ▶ Today lab: practice with `dplyr` verbs (and base R manipulation)
- ▶ Tomorrow lab: more practice in data manipulation

Touchstones: I can comfortably manipulate data²

Next lecture:

- ▶ Using `if` and `ifelse`

²i.e. adjust or add columns to data, subset it in various ways, sort it as needs be and make summary tables.