

Accelerated Lecture 5: Data Visualization

Harris Coding Camp – Accelerate Track

Summer 2022

Today's lesson

- ▶ Why visualize data?
- ▶ How to do so with ggplot
 - ▶ How to map data to aesthetics with `aes()` (and what that means)
 - ▶ How to visualize the mappings with `geoms`
 - ▶ How to get more out of your data by using multiple aesthetics
 - ▶ How to use facets to add dimensionality
- ▶ Some base R tips

We have entire courses on data visualization. This is just a sample.

Data Visualization: Motivation

Suppose we want to know the following info:

- ▶ How have annual housing sales in Texas changed over time?
- ▶ How do these trends compare between cities?

Data Visualization: Motivation


Is simply looking at all these data helpful?

```
## # A tibble: 10 x 2
##   year total_volume
##   <int>         <dbl>
## 1  2000  33342410971
## 2  2001  35804815138
## 3  2002  37798888462
## 4  2003  41674204834
## 5  2004  47913188880
## 6  2005  56534755111
## 7  2006  65237510783
## 8  2007  64393979596
## 9  2008  54198855809
## 10 2009  48450447327
```

Data Visualization: Motivation


What if we make a plot of annual housing sales over time. . .

- ▶ Now we can quickly understand and communicate about our data



Challenge: How do we efficiently communicate
how to visualize data to the computer?

Introducing the “grammar of graphics” and ggplot




- ▶ `gg` = Grammar of graphics (Wickham 2010, Wilkinson et. al 2005, Bertin 1983)
 - ▶ Rules for how to put graph-parts together to make a logical expression
 - ▶ Implemented in R with `ggplot2`, a data visualization package in the tidyverse

Basic Components of ggplot (Layers)

- ▶ Layer 1: Scales and coordinates (`ggplot()`)
 - ▶ A data frame
 - ▶ Aesthetic mapping: how data are mapped to x-axis, y-axis, color, size, etc
- ▶ Layer 2: Geometry (`geom_xxx()`)
 - ▶ geometric objects like points, lines, shapes
- ▶ Layer 3: Labels (`labs()`)
 - ▶ title, legend, etc

What is Aesthetic?

- ▶ An aesthetic is a visual property of the objects in your plot
 - ▶ Including things like size, shape, color or x and y locations
- ▶ To display values, map variables in the data to visual properties of the geom (aesthetics)



Simplest ggplot code structure

```
ggplot(data = dataset,  
       mapping = aes(x = x_variable,  
                     y = y_variable))  
geom_<name>()
```

ggplot() tells R to prepare to make a plot.

```
# Let's prepare new data frame 'annual_sales'  
annual_sales <-  
  txhousing %>%  
  group_by(year) %>%  
  summarize(total_volume = sum(volume, na.rm = TRUE))
```


```
# Layer 1, data frame  
ggplot(data = annual_sales)
```

Layer 1: adding an aesthetic mapping

`mapping = aes()` declares how to map the data to “aesthetics”:

- ▶ R will map each row of the data (`year`, `total_volume`) to the (`x`,`y`)
- ▶ i.e. tell R to make `x`-axis `year` and `y`-axis `total_volume`


```
ggplot(data = annual_sales,  
       mapping = aes(x = year, y = total_volume))
```



Layer 2: visualizing the mapping with geom

Here we see points by using `geom_point()`:

```
ggplot(data = annual_sales,  
       mapping = aes(x = year, y = total_volume)) +  
  geom_point()
```




Layer 2: visualizing the mapping with geom

Here we see bars by using `geom_col()`.

- Each *observation* or row has a (year, total_volume) mapped to the coordinate pair (x,y)


```
ggplot(data = annual_sales,  
       mapping = aes(x = year, y = total_volume)) +  
  geom_col()
```



Layer 2: visualizing the mapping with geom

Here we see a line connecting each (x,y) pair using `geom_line()`.


```
ggplot(data = annual_sales,  
       mapping = aes(x = year, y = total_volume)) +  
  geom_line()
```



Layer 2: visualizing the mapping with geom

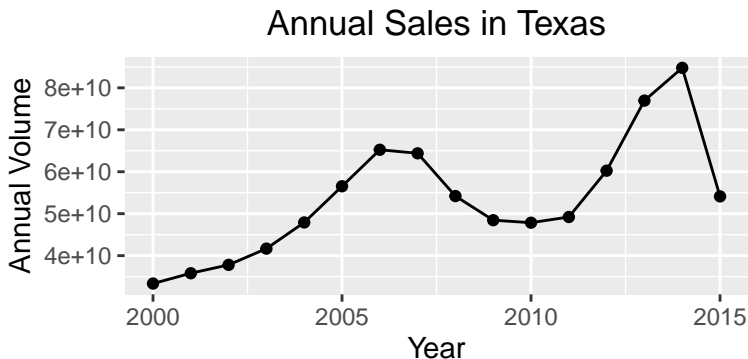
The data can be visualized with different geoms that can be composed (+) together:

```
ggplot(data = annual_sales,  
       mapping = aes(x = year, y = total_volume)) +  
  geom_line() +  
  geom_point()
```




Layer 3: Adding labels makes the plot more readable:

```
ggplot(data = annual_sales,  
       mapping = aes(x = year, y = total_volume)) +  
  geom_line() +  
  geom_point() +  
  labs(x = "Year", y = "Annual Volume",  
       title = "Annual Sales in Texas") +  
  theme(plot.title = element_text(hjust = 0.5)) #center the title
```



Over laying multiple geoms: adding vertical lines

```
annual_sales %>%  
  ggplot(aes(x = year, y = total_volume)) +  
    geom_point() +  
    geom_vline(aes(xintercept = 2007),  
               linetype = "dotted")
```




- ▶ add horizontal lines with `geom_hline()`
- ▶ add any linear fit with `geom_abline()` by providing a slope and intercept

aesthetics beyond the x and y position

We'll use midwest data and start with only mapping to x and y

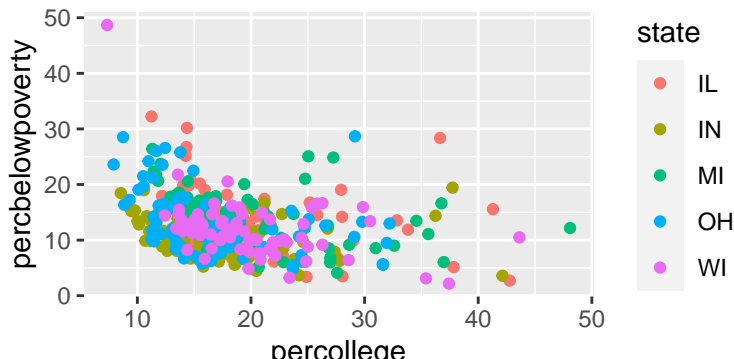
```
midwest %>%  
  ggplot(aes(x = percollege,  
             y = percbelowpoverty)) +  
  geom_point()
```



ggplot(): Using color

- ▶ color maps data to the color of points or lines
 - ▶ Each state is assigned a color
 - ▶ This works with discrete data and continuous data

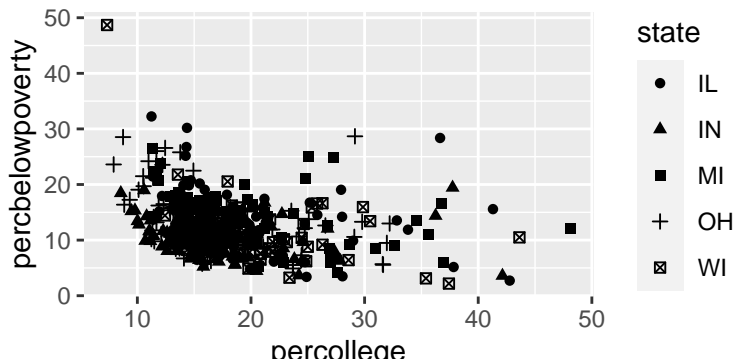
```
midwest %>%  
  ggplot(aes(x = percollege,  
             y = percbelowpoverty,  
             color = state)) +  
  geom_point()
```



ggplot(): Using shape

- ▶ shape maps data to the shape of points
 - ▶ Each state is assigned a shape
 - ▶ This works with discrete data only

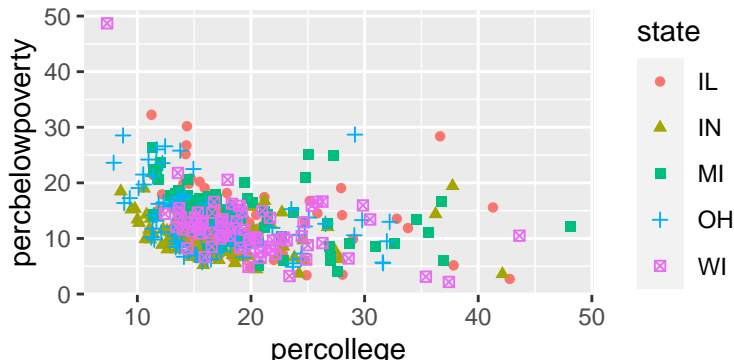
```
midwest %>%  
  ggplot(aes(x = percollege,  
             y = percbelowpoverty,  
             shape = state)) +  
  geom_point()
```



ggplot(): Using color + shape

- ▶ Combining color and shape:
 - ▶ Each state is assigned a shape and color


```
midwest %>%  
  ggplot(aes(x = percollege,  
             y = percbelowpoverty,  
             color = state,  
             shape = state)) +  
  geom_point()
```



ggplot(): Using alpha

- ▶ alpha maps data to the transparency of points
- ▶ we map the percentage of people within a known poverty status to alpha

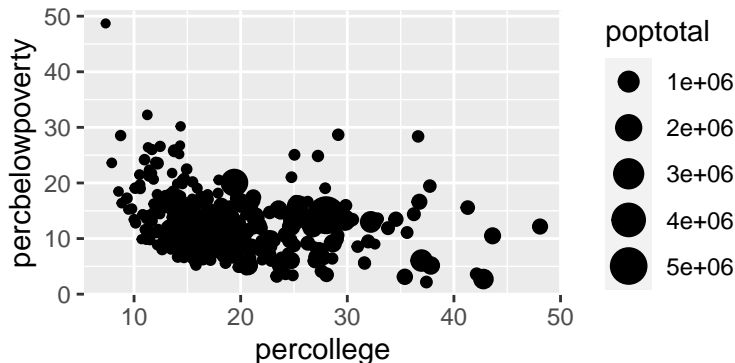
```
midwest %>%  
  ggplot(aes(x = percollege,  
             y = percbelowpoverty,  
             alpha = poptotal)) +  
  geom_point()
```



ggplot(): Using size

- ▶ size maps data to the size of points and width of lines.
- ▶ we map the percentage of people within a known poverty status to size

```
midwest %>%  
  ggplot(aes(x = percollege,  
             y = percbelowpoverty,  
             size = poptotal)) +  
  geom_point()
```




ggplot(): Using multiple aesthetics together

We can combine any and all aesthetics, and even map the same variable to multiple aesthetics

```
midwest %>%  
  ggplot(aes(x = percollege,  
             y = percbelowpoverty,  
             alpha = percpovertyknown,  
             size = poptotal,  
             color = state)) +  
  geom_point()
```


ggplot(): Using multiple aesthetics together



ggplot(): Facets (facet_grid)

Facets provide an additional tool to explore multidimensional data:


```
midwest %>%  
  ggplot(aes(x = log(poptotal),  
             y = percbelowpoverty)) +  
    geom_point() +  
    facet_grid(vars(state))
```



ggplot(): Facets (facet_wrap)

Facets provide an additional tool to explore multidimensional data:

```
midwest %>%  
  ggplot(aes(x = log(poptotal),  
             y = percbelowpoverty)) +  
    geom_point() +  
    facet_wrap(vars(state))
```



ggplot(): Using aesthetics to explore data

Different geoms have specific aesthetics that go with them.

- ▶ the ggplot cheatsheet shows all the geoms with their associated aesthetics

Data visualization with ggplot2 : CHEAT SHEET



Basics

ggplot2 is based on the **grammar of graphics**, the idea that you can build every graph from the same components: a **data set**, a **coordinate system**, and **geoms**—visual marks that represent data points.



To display values, map variables in the data to visual properties of the geom (**aesthetics**) like **size**, **color**, and **x** and **y** locations.



Complete the template below to build a graph.

```
ggplot(data = <DATA>) +  
  <GEOM_FUNCTION>(mapping = aes(<MAPPINGS>))  
stat = <STAT>, position = <POSITION> +  
<COORDINATE_FUNCTION> +  
<FACET_FUNCTION> +  
<SCALE_FUNCTION> +  
<THEME_FUNCTION>
```

Not required, sensible defaults supplied

ggplot(data = mpg, aes(x = cty, y = hwy)) Begins a plot that you finish by adding layers to. Add one geom function per layer.

last_plot() Returns the last plot.

ggsave("plot.png", width = 5, height = 5) Saves last plot as 5 x 5 file named "plot.png" in working directory. Matches file type to file extension.

Geoms

Use a geom function to represent data points, use the geom's aesthetic properties to represent variables. Each function returns a layer.

GRAPHICAL PRIMITIVES

a <- ggplot(economics, aes(date, unemploy))
b <- ggplot(seals, aes(x = long, y = lat))

- a + geom_blank()** and **a + expand_limits()**
Ensure limits include values across all plots.
- b + geom_curve**(aes(yend = lat + 1, xend = long + 1), curvature = 1, x, yend, y, yend, alpha, angle, color, fill, size, stroke, size)
- a + geom_path**(lineend = "butt", linejoin = "round", linemitre = 1)
x, y, alpha, color, group, linetype, size
- a + geom_polygon**(aes(alpha = 50)) - x, y, alpha, color, fill, group, subgroup, linetype, size
- b + geom_rect**(aes(xmin = long, ymin = lat, xmax = long + 1, ymax = lat + 1)) - xmax, xmin, ymax, ymin, alpha, color, fill, linetype, size
- a + geom_ribbon**(aes(ymin = unemploy - 900, ymax = unemploy + 900)) - x, ymax, ymin, alpha, color, fill, group, linetype, size

LINE SEGMENTS

- common aesthetics: x, y, alpha, color, linetype, size
- b + geom_abline**(aes(intercept = 0, slope = 1))
- b + geom_hline**(aes(intercept = lat))
- b + geom_vline**(aes(intercept = long))
- b + geom_segment**(aes(yend = lat + 1, xend = long + 1))
- b + geom_spoke**(aes(angle = 1:1155, radius = 1))

ONE VARIABLE continuous

- c <- ggplot(mpg, aes(hwy)); c2 <- ggplot(mpg)**
- c + geom_area**(stat = "bin")
x, y, alpha, color, fill, linetype, size
- c + geom_density**(kernel = "gaussian")
x, y, alpha, color, fill, group, linetype, size, weight

TWO VARIABLES both continuous

- e <- ggplot(mpg, aes(cty, hwy))**
- e + geom_label**(aes(label = cty, nudges_x = 1, nudges_y = 1), x, y, label, alpha, angle, color, family, fontface, hjust, lineheight, size, vjust)
- e + geom_point**
x, y, alpha, color, fill, shape, size, stroke, size
- e + geom_quantile**
x, y, alpha, color, group, linetype, size, weight
- e + geom_rug**(sides = "b")
x, y, alpha, color, linetype, size
- e + geom_smooth**(method = lm)
x, y, alpha, color, fill, group, linetype, size, weight
- e + geom_text**(aes(label = cty), nudges_x = 1, nudges_y = 1) - x, y, label, alpha, angle, color, family, fontface, hjust, lineheight, size, vjust

one discrete, one continuous

- f <- ggplot(mpg, aes(class, hwy))**
- f + geom_col**
x, y, alpha, color, fill, group, linetype, size
- f + geom_boxplot**
x, y, lower, middle, upper, ymax, ymin, alpha, color, fill, group, linetype, shape, size, weight
- f + geom_dotplot**(binaxis = "y", stackdir = "center")
x, y, alpha, color, fill, group
- f + geom_violin**(scale = "area")
x, y, alpha, color, fill, group, linetype, size, weight

both discrete

continuous bivariate distribution

- h <- ggplot(diamonds, aes(carat, price))**
- h + geom_bin2d**(binwidth = c(0.25, 500))
x, y, alpha, color, fill, linetype, size, weight
- h + geom_density_2d()**
x, y, alpha, color, group, linetype, size
- h + geom_hex()**
x, y, alpha, color, fill, size

continuous function

- i <- ggplot(economics, aes(date, unemploy))**
- i + geom_area**
x, y, alpha, color, fill, linetype, size
- i + geom_line**
x, y, alpha, color, group, linetype, size
- i + geom_step**(direction = "hv")
x, y, alpha, color, group, linetype, size


visualizing error

- df <- data.frame(gp = c("A", "B"), fit = 4:5, se = 1:2)**
- j <- ggplot(df, aes(gp, fit, ymin = fit - se, ymax = fit + se))**
- j + geom_crossbar**(fatten = 2) - x, y, ymax, ymin, alpha, color, fill, group, linetype, size
- j + geom_errorbar**(- x, ymax, ymin, alpha, color, group, linetype, size, width)
- Also geom_errorbarh()**
- j + geom_linerange**
x, ymin, ymax, alpha, color, group, linetype, size
- j + geom_pointrange**(- x, y, ymin, ymax, alpha, color, fill, group, linetype, shape, size)

maps

Try it yourself: Plot 1

1. Adjust code to reproduce the following plot (sample codes provided in the next slide):

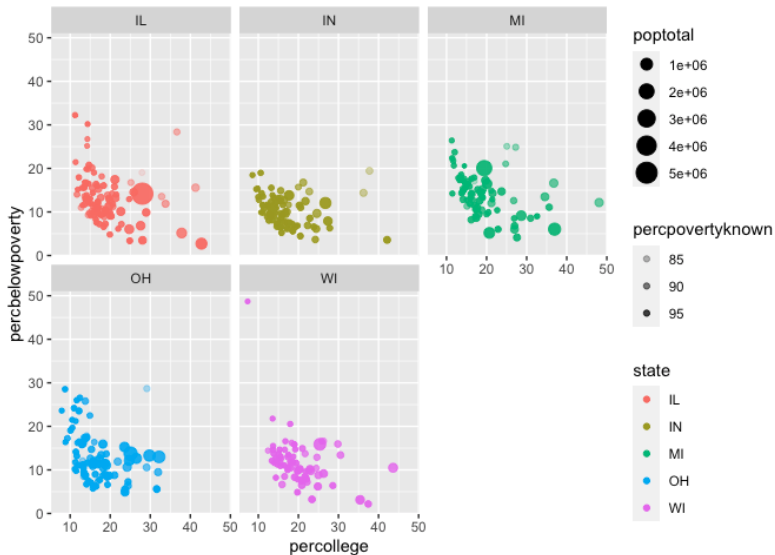


Try it yourself: Plot 1

```
midwest %>%  
  ggplot(aes(x = ?,  
             y = ?,  
             )) +  
  geom_?()
```

Try it yourself: Plot 2

- Adjust code to reproduce the following plot (sample codes provided in the next slide):



Try it yourself: Plot 2

```
midwest %>%  
  ggplot(aes(x = ?,  
             y = ?,  
             color = state,  
             size = ?,  
             alpha = percpovertyknown)) +  
  geom_?() +  
  facet_wrap(vars(?))
```

discrete vs continuous data


aes	discrete	continuous
	limited number of classes usually chr or lgl	unlimited number of classes numeric
x, y	yes	yes
color, fill	yes	yes
shape	yes (6 or fewer categories)	no
size, alpha	not advised	yes
facet	yes	not advised

Here, discrete and continuous have different meaning than in math

- ▶ For ggplot meaning is more fluid.
- ▶ If there are fewer than 6 to 10 groups, discrete visualizations can work
- ▶ If your “discrete” data is numeric, use `as.character()` or `as_factor()` to enforce the decision.

color can be continuous


```
midwest %>%  
  ggplot(aes(x = percollege,  
             y = percbelowpoverty,  
             color = percpovertyknown)) +  
  geom_point()
```



shape does not play well with many categories

- ▶ Will only map to 6 categories, the rest become NA.
- ▶ We can override this behavior and get up to 25 distinct shapes


```
midwest %>%  
  ggplot(aes(x = percollege,  
             y = percbelowpoverty,  
             shape = county)) +  
  geom_point() +  
  # legend off, otherwise it overwhelms  
  theme(legend.position = "none")
```



alpha and size can be misleading with discrete data

```
midwest %>%  
  ggplot(aes(x = percollege,  
             y = percbelowpoverty,  
             alpha = state)) +  
  geom_point()
```

Warning: Using alpha for a discrete variable is not advised



Type of figures

1. Distribution of **univariate (single variable)**

- ▶ bar plot, histogram, density plot, etc

2. Relationship between **bivariate (two variables)**


- ▶ scatter plot, line plot, boxplot, (segmented) bar plot, etc

3. Relationship between **many variables** at once

- ▶ usually focusing on the relationship between two while conditioning for others


Univariate: bar plot

```
midwest %>%  
  ggplot(aes(x = state)) +  
    geom_bar()
```



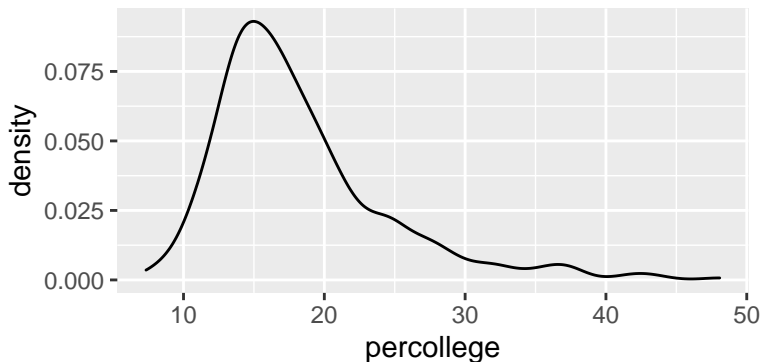
Univariate: histogram

```
midwest %>%  
  ggplot(aes(x = percollege)) +  
    geom_histogram(binwidth = 1)
```




Univariate: density

```
midwest %>%  
  ggplot(aes(x = percollege)) +  
    geom_density()
```




Univariate: box plots

```
midwest %>%  
  ggplot(aes(x = percollege)) +  
    geom_boxplot()
```




Bivariate: scatter plot

```
midwest %>%  
  ggplot(aes(x = percollege,  
             y = percbelowpoverty)) +  
  geom_point(size=1)
```




Bivariate: scatter + smooth Line plot

```
midwest %>%  
  ggplot(aes(x = percollege,  
             y = percbelowpoverty)) +  
    geom_point(size = 1) +  
    geom_smooth()
```




Bivariate: scatter + smooth Line plot

```
midwest %>%  
  ggplot(aes(x = percollege,  
             y = percbelowpoverty)) +  
    geom_point(size = 1) +  
    geom_smooth(se = FALSE) # turn off std errors
```



Bivariate: box plots

```
midwest %>%  
  ggplot(aes(x = state,  
             y = percollege)) +  
    geom_boxplot()
```



Recap

- ▶ Visualizing our data can help lead to powerful insights between variable relationships
 - ▶ Making quick plots helps us understand data and makes us aware of data issues
- ▶ `ggplot` starts by mapping data to “aesthetics”
 - ▶ e.g. What data shows up on x and y axes and how color, size and shape appear on the plot
- ▶ Then, we use `geoms` to create a visualization based on the mapping
- ▶ We many consider adding labels to make plots more readable
- ▶ There are many ways you can visualize your data!

Next steps

Labs

- ▶ Today: Data visualization with `ggplot` (may run into tomorrow)
- ▶ Tomorrow: Introducing plotting in base R

I can produce basic plots to explore and communicate about data

Lecture

- ▶ Data manipulation and analysis with groups

Appendix: Some graphs you made along the way


- ▶ Distributions
- ▶ Grouped bar graph
- ▶ Faceted bar graph

Appendix: distributions

- ▶ `geom_density()` only requires an `x` aesthetic and it calculates the distribution to plot.
- ▶ We can set the aesthetics manually, independent of data for nicer graphs.

```
chi_sq_samples <-  
  tibble(x = c(rchisq(100000, 2),  
               rchisq(100000, 3),  
               rchisq(100000, 4)),  
         df = rep(c("2", "3", "4"), each = 1e5))  
  
chi_sq_samples %>%  
  ggplot(aes(x = x, fill = df)) +  
  geom_density(alpha = .5) +  
  labs(fill = "df", x = "sample")
```

Appendix: distributions




Appendix: grouped bar graph

- ▶ `position = "dodge2"` tells R to put bars next to each other, rather than stacked on top of each other.
- ▶ Notice we use `fill` and not `color` because we're "filling" an area.

```
mean_share_per_country %>%  
  ggplot(aes(y = country,  
             x = mean_share,  
             fill = percentile)) +  
  geom_col(position = "dodge2") +  
  labs(x = "Mean share of national wealth",  
       y = "",  
       fill = "Wealth\npercentile")
```

Appendix: grouped bar graph



Appendix: faceted bar graph

- ▶ Notice that we manipulate our data to the right specification before making this graph
- ▶ Using `facet_wrap` we get a distinct graph for each time period.

```
mean_share_per_country_with_time %>%  
  ggplot(aes(x = country,  
             y = mean_share,  
             fill = percentile)) +  
  geom_col(position = "dodge2") +  
  facet_wrap(vars(time_period))
```

Appendix: faceted bar graph

