

Accelerated Lecture 6: Grouped Analysis

Harris Coding Camp

Summer 2022

Grouping data with `group_by()`

Often we want to repeat the same analysis across different subgroups. We can automate that with `group_by()`. We will:

- ▶ summarize by group with `group_by() + summarize()`
- ▶ create new columns with `group_by() + mutate()`
- ▶ `filter()` data with group specific matching criteria

Grouping data with group_by()

By itself, group_by() doesn't do much. But, once an object is grouped, **all subsequent dplyr functions are run separately “by group”**:

- ▶ count() returns the number of observations

```
txhousing %>% count()
#> # A tibble: 1 x 1
#>       n
#>   <int>
#> 1  8602
txhousing %>% group_by(year) %>% count() %>% glimpse()
#> Rows: 16
#> Columns: 2
#> Groups: year [16]
#> $ year <int> 2000, 2001, 2002, 2003, 2004, 2005, 2006, 2007, 2008, 2009, 2010, 2011, 2012, 2013, 2014, 2015
#> $ n      <int> 552, 552, 552, 552, 552, 552, 552, 552, 552, 552, 552, 552, 552, 552, 552, 552
```

Grouping data with group_by()

Ungrouped data ...

```
txhousing %>% glimpse()
#> Rows: 8,602
#> Columns: 9
#> $ city      <chr> "Abilene", "Abilene", "Abilene", "Abi
#> $ year      <int> 2000, 2000, 2000, 2000, 2000, 2000, 20
#> $ month     <int> 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12
#> $ sales     <dbl> 72, 98, 130, 98, 141, 156, 152, 131,
#> $ volume    <dbl> 5380000, 6505000, 9285000, 9730000, 10
#> $ median    <dbl> 71400, 58700, 58100, 68600, 67300, 663
#> $ listings  <dbl> 701, 746, 784, 785, 794, 780, 742, 763
#> $ inventory <dbl> 6.3, 6.6, 6.8, 6.9, 6.8, 6.6, 6.2, 6.2
#> $ date      <dbl> 2000.000, 2000.083, 2000.167, 2000.250
```

Grouping data with group_by()

group_by() adds meta-data indicating which rows/observations belong to which group.

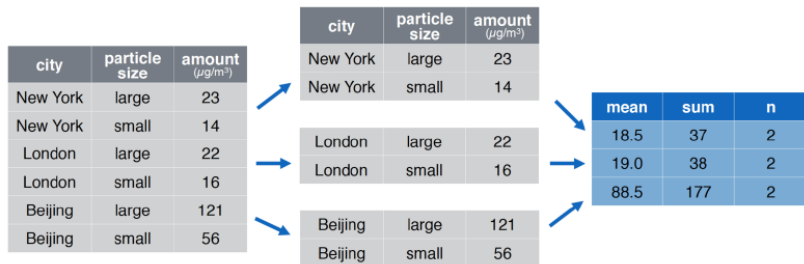
```
group_by(txhousing, city) %>% glimpse()
#> Rows: 8,602
#> Columns: 9
#> Groups: city [46]
#> $ city      <chr> "Abilene", "Abilene", "Abilene", "Abi
#> $ year      <int> 2000, 2000, 2000, 2000, 2000, 2000, 20
#> $ month     <int> 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12
#> $ sales     <dbl> 72, 98, 130, 98, 141, 156, 152, 131,
#> $ volume    <dbl> 5380000, 6505000, 9285000, 9730000, 10
#> $ median    <dbl> 71400, 58700, 58100, 68600, 67300, 66
#> $ listings  <dbl> 701, 746, 784, 785, 794, 780, 742, 76
#> $ inventory <dbl> 6.3, 6.6, 6.8, 6.9, 6.8, 6.6, 6.2, 6.
#> $ date      <dbl> 2000.000, 2000.083, 2000.167, 2000.250
```

We can group by multiple columns

We now have $46 \text{ cities} \times 16 \text{ years} = 736 \text{ groups!}$

```
group_by(txhousing, city, year) %>% glimpse()
#> Rows: 8,602
#> Columns: 9
#> Groups: city, year [736]
#> $ city      <chr> "Abilene", "Abilene", "Abilene", "Abi
#> $ year      <int> 2000, 2000, 2000, 2000, 2000, 2000, 20
#> $ month     <int> 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12
#> $ sales     <dbl> 72, 98, 130, 98, 141, 156, 152, 131,
#> $ volume    <dbl> 5380000, 6505000, 9285000, 9730000, 10
#> $ median    <dbl> 71400, 58700, 58100, 68600, 67300, 66
#> $ listings  <dbl> 701, 746, 784, 785, 794, 780, 742, 76
#> $ inventory <dbl> 6.3, 6.6, 6.8, 6.9, 6.8, 6.6, 6.2, 6.
#> $ date      <dbl> 2000.000, 2000.083, 2000.167, 2000.250
```

Grouped summary with `group_by()` `%>% summarize()`



In this example, we want to calculate the mean and sum of amount by city

- ▶ we must repeat the same analysis across different groups (i.e. cities)
- ▶ Using `group_by()` `%>% summarize()` makes it a lot easier!

Grouped summary with group_by() %>% summarize()

Let's see the upgrouped summary statistics first:

```
txhousing %>%  
  summarize(total_sales = sum(sales, na.rm = TRUE),  
            total_volume = sum(volume, na.rm = TRUE),  
            mean_house_price = total_volume / total_sales)  
#> # A tibble: 1 x 3  
#>   total_sales total_volume mean_house_price  
#>   <dbl>         <dbl>         <dbl>  
#> 1    4415202 858502159353    194442.
```


Grouped summary with `group_by()` `%>% summarize()`

Use case: You want summary statistics for all subsets (i.e. each year):

- ▶ `summarize()` collapses a data frame to a single row for each group
- ▶ The count function `n()` takes no arguments and returns the size of a group

```
annual_housing_prices <-  
  txhousing %>%  
  group_by(year) %>%  
  summarize(n_within_group = n(),  
            total_sales = sum(sales, na.rm = TRUE),  
            total_volume = sum(volume, na.rm = TRUE),  
            mean_house_price = total_volume / total_sales)
```


```
head(annual_housing_prices, n=5)
```

```
#> # A tibble: 5 x 5
```

```
#>   year n_within_group total_sales total_volume mean_house_price  
#>   <int>         <int>         <dbl>         <dbl>         <dbl>  
#> 1  2000             552      222483    33342410971      149865.  
#> 2  2001             552      231453    35804815138      154696.  
#> 3  2002             552      234600    37798888462      161121.  
#> 4  2003             552      253909    41674204834      164130.  
#> 5  2004             552      283999    47913188880      168709
```

How have Texas housing prices changed over time?

```
annual_housing_prices %>%  
  ggplot(aes(x = year,  
             y = mean_house_price)) +  
  geom_line()
```



Grouped summary with `group_by()` + `summarize()`

Use case: You want summary statistics for certain subsets (each year) in a specific city (e.g. Houston):

```
txhousing %>%  
  filter(city == "Houston") %>%  
  group_by(year) %>%  
  summarize(n_within_group = n(),  
            total_sales = sum(sales, na.rm = TRUE),  
            total_volume = sum(volume, na.rm = TRUE),  
            mean_house_price = total_volume / total_sales)
```

#> # A tibble: 16 x 5

#>	year	n_within_group	total_sales	total_volume	mean_house_price
#>	<int>	<int>	<dbl>	<dbl>	<dbl>
#>	1 2000	12	52459	8041166317	153285.
#>	2 2001	12	53856	8541022943	158590.
#>	3 2002	12	56563	9486396667	167714.
#>	4 2003	12	60732	10417774768	171537.
#>	5 2004	12	66979	11776381072	175822.
#>	6 2005	12	72800	13504202605	185497.
#>	7 2006	12	80994	15816104590	195275.
#>	8 2007	12	77668	15789736644	203298.
#>	9 2008	12	65169	13396719487	205569.
#>	10 2009	12	60106	12035965014	200246.

Grouped summary with `group_by()` + `summarize()`

Use case: You want summary statistics for certain subsets (each year, each city) of the data.

```
annual_city_housing_prices <-  
  txhousing %>%  
    group_by(city, year) %>%  
    summarize(total_sales = sum(sales, na.rm = TRUE),  
              total_volume = sum(volume, na.rm = TRUE),  
              mean_house_price = total_volume / total_sales)
```

```
head(annual_city_housing_prices, n=5)
```

```
#> # A tibble: 5 x 5
```

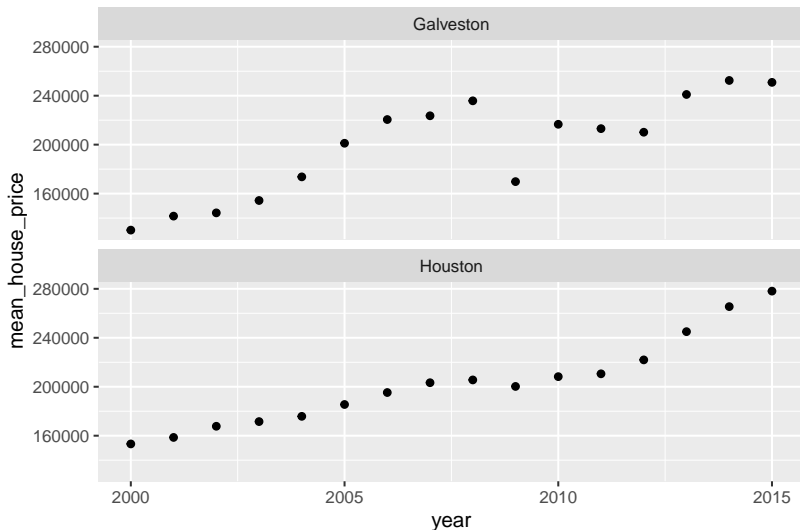
```
#> # Groups:   city [1]
```

#>	city	year	total_sales	total_volume	mean_house_price
#>	<chr>	<int>	<dbl>	<dbl>	<dbl>
#> 1	Abilene	2000	1375	108575000	78964.
#> 2	Abilene	2001	1431	114365000	79920.
#> 3	Abilene	2002	1516	118675000	78282.
#> 4	Abilene	2003	1632	135675000	83134.
#> 5	Abilene	2004	1830	159670000	87251.

How have Texas housing prices changed over time in certain cities?

```
txhousing %>%  
  group_by(city, year) %>%  
  summarize(total_sales = sum(sales, na.rm = TRUE),  
            total_volume = sum(volume, na.rm = TRUE),  
            mean_house_price = total_volume / total_sales) %>%  
  filter(city %in% c("Houston", "Galveston")) %>%  
  ggplot(aes(x = year, y = mean_house_price)) +  
    geom_point() +  
    facet_wrap(facets = "city", nrow = 2)
```

How have Texas housing prices changed over time in certain cities?



Grouping + Summarizing: Base R vs Tidyverse

Base R:

```
# use formula to indicate grouping vars
aggregate(formula = sales ~ city + year,
           data = txhousing,
           FUN = mean)
aggregate(formula = sales ~ city + year,
           data = txhousing,
           FUN = sd)
```

Tidyverse:

```
txhousing %>%
  group_by(city, year) %>%
  summarize(mean_sales = mean(sales, na.rm = TRUE),
            sd_sales = sd(sales, na.rm = TRUE))
```

Ungrouping data

To get rid of groups, use `ungroup()`

```
txhousing_grouped <- group_by(txhousing, year)
```

```
txhousing_grouped %>%  
  ungroup() %>%  
  summarize(total_sales = sum(sales, na.rm = TRUE))  
#> # A tibble: 1 x 1  
#>   total_sales  
#>       <dbl>  
#> 1      4415202
```


What's going on here?

```
txhousing_grouped %>%  
  select(-year) %>%  
  head()
```

```
#> # A tibble: 6 x 9
```

```
#> # Groups:   year [1]
```

```
#>   year city      month sales    volume median listings in  
#>   <int> <chr>    <int> <dbl>    <dbl>    <dbl>    <dbl>  
#> 1  2000 Abilene      1     72  5380000  71400     701  
#> 2  2000 Abilene      2     98  6505000  58700     746  
#> 3  2000 Abilene      3    130  9285000  58100     784  
#> 4  2000 Abilene      4     98  9730000  68600     785  
#> 5  2000 Abilene      5    141 10590000  67300     794  
#> 6  2000 Abilene      6    156 13910000  66900     780
```

grouped data require the grouping variable

We got the message:

Adding missing grouping variables: 'year'

```
txhousing_grouped %>%  
  ungroup() %>%  
  select(-year) %>%  
  head()
```

```
#> # A tibble: 6 x 8
```

```
#>   city      month sales    volume median listings inventory  
#>   <chr>    <int> <dbl>    <dbl>   <dbl>   <dbl>    <dbl>  
#> 1 Abilene      1    72  5380000  71400    701    6.3  
#> 2 Abilene      2    98  6505000  58700    746    6.6  
#> 3 Abilene      3   130  9285000  58100    784    6.8  
#> 4 Abilene      4    98  9730000  68600    785    6.9  
#> 5 Abilene      5   141 10590000  67300    794    6.8  
#> 6 Abilene      6   156 13910000  66900    780    6.6
```

Try it yourself

1. Using `txhousing`, filter observations where `city` is Brazoria County and group by year.
 2. Next, create a variable that represents the total sales in each year.
 3. Plot the total sales over time.
 4. Create two variables that represent the sum of missing & non-missing obs for sales in Brazoria County.
- Whenever you do any aggregation, it's always a good idea to include either a count of missing values `sum(is.na(x))` or a count of non-missing values `sum(!is.na(x))`. That way, you can check that you are not making conclusions based on very small amounts of data!

group_by() %>% mutate() creates new columns with groups in mind

Often we want the “summary” information in the data context.

```
ex <- tibble(period = c(1, 2, 1, 2),  
              group = c("a", "a", "b", "b"),  
              x = c(3, 1, 11, 13),)
```

```
ex %>%  
  group_by(group) %>%  
  mutate(group_mean = mean(x))  
#> # A tibble: 4 x 4  
#> # Groups:   group [2]  
#>   period group      x group_mean  
#>   <dbl> <chr> <dbl>      <dbl>  
#> 1     1 a         3         2  
#> 2     2 a         1         2  
#> 3     1 b        11        12  
#> 4     2 b        13        12
```

Grouped mutate: differences

Use case: You want to work with differences.

```
ex %>%  
  group_by(group) %>%  
  mutate(x_lag = lag(x),  
         x_diff = x - lag(x))  
#> # A tibble: 4 x 5  
#> # Groups:   group [2]  
#>   period group      x x_lag x_diff  
#>   <dbl> <chr> <dbl> <dbl> <dbl>  
#> 1     1 a      3    NA    NA  
#> 2     2 a      1     3    -2  
#> 3     1 b     11    NA    NA  
#> 4     2 b     13    11     2
```

Why is this wrong?

```
ex %>%  
  mutate(x_lag = lag(x),  
         x_diff = x - lag(x))  
#> # A tibble: 4 x 5  
#>   period group      x x_lag x_diff  
#>   <dbl> <chr> <dbl> <dbl> <dbl>  
#> 1     1 a      3    NA    NA  
#> 2     2 a      1     3    -2  
#> 3     1 b     11     1    10  
#> 4     2 b     13    11     2
```

Grouped mutate: differences

Use case: You want to work with differences. (Try running the code without `group_by()` and carefully compare the results.)

```
july_texas_txhousing <-  
  txhousing %>%  
    filter(month == 7) %>%  
    select(city, year, sales)  
  
differenced_data <-  
  july_texas_txhousing %>%  
    group_by(city) %>%  
    mutate(last_year_sales = lag(sales),  
           diff_sales = sales - lag(sales))
```

Grouped mutate: differences

Use case: You want to work with differences between sales in different years.¹

```
differenced_data %>% head(5)
#> # A tibble: 5 x 5
#> # Groups:   city [1]
#>   city      year sales last_year_sales diff_sales
#>   <chr>   <int> <dbl>         <dbl>         <dbl>
#> 1 Abilene  2000    152          NA           NA
#> 2 Abilene  2001    134         152          -18
#> 3 Abilene  2002    159         134           25
#> 4 Abilene  2003    171         159           12
#> 5 Abilene  2004    176         171            5
```

¹lag()'s sibling is lead() which will give you data from the following year.

Grouped mutate: other window functions

- ▶ See the “Data transformation with dplyr” cheatsheet (page 2) for more vectorized window functions.

```
ex %>%
  group_by(group) %>%
  mutate(cumulative = cumsum(x),
         # comparing values to summaries
         centered = (x - mean(x)))
```

#> # A tibble: 4 x 5

#> # Groups: group [2]

#>	period	group	x	cumulative	centered
#>	<dbl>	<chr>	<dbl>	<dbl>	<dbl>
#> 1	1	a	3	3	1
#> 2	2	a	1	4	-1
#> 3	1	b	11	11	-1
#> 4	2	b	13	24	1

Grouped mutate: ranking

```
ex %>%  
  group_by(group) %>%  
  mutate(rank = rank(desc(x))) %>%  
  arrange(group, rank)  
#> # A tibble: 4 x 4  
#> # Groups:   group [2]  
#>   period group      x  rank  
#>   <dbl> <chr> <dbl> <dbl>  
#> 1       1 a         3     1  
#> 2       2 a         1     2  
#> 3       2 b        13     1  
#> 4       1 b        11     2
```

Aside: Grouped arrange

Say you want to order the data without explicitly adding a rank.

- ▶ Almost all dplyr function will operate group-by-group on grouped data.
- ▶ `arrange()` is an exception

```
ex %>%  
  group_by(group) %>%  
  arrange(row_number(desc(x)))
```

```
#> # A tibble: 4 x 3  
#> # Groups:   group [2]  
#>   period group      x  
#>   <dbl> <chr> <dbl>  
#> 1       2 b      13  
#> 2       1 b      11  
#> 3       1 a       3  
#> 4       2 a       1
```

Aside: Grouped arrange

However, you only need to add the grouping column(s) or `.by_group = TRUE` to get the desired output.

```
ex %>%  
  arrange(group, row_number(desc(x)))
```

```
ex %>%  
  group_by(group) %>%  
    # this option is nice if you have many grouping cols  
    arrange(row_number(desc(x)), .by_group = TRUE)  
#> # A tibble: 4 x 3  
#> # Groups:   group [2]  
#>   period group      x  
#>   <dbl> <chr> <dbl>  
#> 1       1 a         3  
#> 2       2 a         1  
#> 3       2 b        13  
#> 4       1 b        11
```

Grouped mutate: ranking

Use case: You want to rank sales within group. (Try running the code without `group_by()` and carefully compare the results.)

```
ranked_data <-  
july_texas_txhousing %>%  
  group_by(year) %>%  
  mutate(sales_rank = rank(desc(sales)))
```

Grouped mutate: ranking

Use case: You want to rank sales within group.²

```
ranked_data %>% arrange(year, sales_rank) %>% head(10)
```

#> # A tibble: 10 x 4

#> # Groups: year [1]

#>	city	year	sales	sales_rank
#>	<chr>	<int>	<dbl>	<dbl>
#>	1 Houston	2000	5009	1
#>	2 Dallas	2000	4276	2
#>	3 Austin	2000	1818	3
#>	4 San Antonio	2000	1508	4
#>	5 Collin County	2000	1007	5
#>	6 Fort Bend	2000	753	6
#>	7 NE Tarrant County	2000	686	7
#>	8 Denton County	2000	638	8
#>	9 Fort Worth	2000	548	9
#>	10 Montgomery County	2000	463	10

²R has a variety of related functions see ?ranking

Grouped filter

Use case: You want to work with the top 5 cities for each year.

```
ranked_data %>%  
  # we already added ranks!  
  filter(sales_rank <= 5) %>%  
  arrange(year, sales_rank) %>%  
  head()  
#> # A tibble: 6 x 4  
#> # Groups:   year [2]  
#>   city          year sales sales_rank  
#>   <chr>         <int> <dbl>     <dbl>  
#> 1 Houston      2000  5009         1  
#> 2 Dallas        2000  4276         2  
#> 3 Austin        2000  1818         3  
#> 4 San Antonio   2000  1508         4  
#> 5 Collin County 2000  1007         5  
#> 6 Houston      2001  5424         1
```

Grouped filter

Use case: You want to work with the top 5 cities for each year.

```
july_texas_txhousing %>%  
  group_by(year) %>%  
  # we don't need sales_rank to filter by ranks!  
  filter(rank(desc(sales)) <= 5) %>%  
  arrange(year, desc(sales)) %>%  
  head()  
  
#> # A tibble: 6 x 3  
#> # Groups:   year [2]  
#>   city          year sales  
#>   <chr>         <int> <dbl>  
#> 1 Houston      2000  5009  
#> 2 Dallas       2000  4276  
#> 3 Austin       2000  1818  
#> 4 San Antonio  2000  1508  
#> 5 Collin County 2000  1007  
#> 6 Houston     2001  5424
```


count() is a useful short cut

Based on what you know about txhousing. Can you tell what count() does?

```
txhousing %>%  
  count(city, year) %>%  
  head(5)
```

#> # A tibble: 5 x 3

<i>#></i>	<i>city</i>	<i>year</i>	<i>n</i>
<i>#></i>	<i><chr></i>	<i><int></i>	<i><int></i>
<i>#> 1</i>	<i>Abilene</i>	<i>2000</i>	<i>12</i>
<i>#> 2</i>	<i>Abilene</i>	<i>2001</i>	<i>12</i>
<i>#> 3</i>	<i>Abilene</i>	<i>2002</i>	<i>12</i>
<i>#> 4</i>	<i>Abilene</i>	<i>2003</i>	<i>12</i>
<i>#> 5</i>	<i>Abilene</i>	<i>2004</i>	<i>12</i>

count() is a useful short cut

count() does the following:

- ▶ applies group_by() on the specified column(s)
- ▶ applies summarize() and returns column n with the number of rows per group
- ▶ applies ungroup()

```
txhousing %>%  
  group_by(city, year) %>%  
  summarize(n = n()) %>%  
  ungroup() %>%  
  head(5)
```

```
#> # A tibble: 5 x 3  
#>   city      year      n  
#>   <chr>   <int> <int>  
#> 1 Abilene  2000    12  
#> 2 Abilene  2001    12  
#> 3 Abilene  2002    12  
#> 4 Abilene  2003    12  
#> 5 Abilene  2004    12
```

add_count() is a useful short cut

Based on what you know about txhousing. Can you tell what add_count() does?

```
txhousing %>%  
  select(city, year, sales) %>%  
  add_count(city, year) %>%  
  head(5)
```

```
#> # A tibble: 5 x 4  
#>   city      year sales      n  
#>   <chr>    <int> <dbl> <int>  
#> 1 Abilene  2000     72     12  
#> 2 Abilene  2000     98     12  
#> 3 Abilene  2000    130     12  
#> 4 Abilene  2000     98     12  
#> 5 Abilene  2000    141     12
```

add_count() is a useful short cut

add_count() does the following:

- ▶ applies group_by() on the specified column(s)
- ▶ applies mutate() to add a new column n with the counts of rows per group while retaining all the other data frame columns
- ▶ applies ungroup()

```
txhousing %>%  
  select(city, year, sales) %>%  
  group_by(city, year) %>%  
  mutate(n = n()) %>%  
  ungroup() %>%  
  head(5)
```

```
#> # A tibble: 5 x 4  
#>   city      year sales      n  
#>   <chr>   <int> <dbl> <int>  
#> 1 Abilene  2000     72     12  
#> 2 Abilene  2000     98     12  
#> 3 Abilene  2000    130     12  
#> 4 Abilene  2000     98     12  
#> 5 Abilene  2000    141     12
```

Try it yourself: Setup

midwest is a data set that comes bundled with tidyverse. First, let's calculate the total population of Ohio in the following way:

```
midwest %>% filter(state == "OH") %>%  
  summarize(total_population = sum(poptotal))  
#> # A tibble: 1 x 1  
#>   total_population  
#>           <int>  
#> 1       10847115
```

With group_by, you can calculate the total population of all the states at once!

```
midwest %>% group_by(state) %>%  
  summarize(total_population = sum(poptotal))  
#> # A tibble: 5 x 2  
#>   state total_population  
#>   <chr>           <int>  
#> 1 IL             11430602  
#> 2 IN             5544159  
#> 3 MI             9295297  
#> 4 OH             10847115  
#> 5 WI             4891769
```

Try it yourself: `group_by()`

1. For each state in the midwest data, calculate the total area.
2. For each state in the midwest data, calculate the proportion of counties that are in a metro area (`inmetro`).³
3. Add a column to midwest called `pop_state` that equals the state population. Compare your result to what you calculated early.

```
# fill in the ... with appropriate code
midwest %>%
  group_by( ... ) %>%
  mutate(pop_state = ... )
```

4. Building off the previous question, create a column that shows the number of people living below the poverty line (`percbelowpoverty`) in each county.

³Recall that `mean()` of a column of 0 and 1s tell you the proportion of 1s.

Try it yourself: `count()`

5. Reproduce this table using `count()`.

```
#> # A tibble: 2 x 2
#>   inmetro     n
#>   <int> <int>
#> 1     0  287
#> 2     1  150
```

Recap: Analysis by group with dplyr

This lesson gave you an idea about how to:

- ▶ summarize data by group with `group_by()` + `summarize()`
- ▶ created new columns with `group_by()` + `mutate()`
 - ▶ we saw `lag()` and `rank()`, but you could get also add group-level stats like `mean()` and `median()`
- ▶ `filter()` data with group specific matching criteria
- ▶ use `count()` and `add_count()` as short cuts for getting group level counts

Next steps:

Lab:

- ▶ Today: Grouped analysis

I can streamline analysis of subgroup data using `group_by()` and `dplyr` verbs

Lecture:

- ▶ Tomorrow: Writing your own functions!